



UNIVERSITÀ  
di **VERONA**

Dipartimento  
di **INFORMATICA**



---

**OpenMP™**

Part I

*Laurea magistrale in Ingegneria e Scienze Informatiche  
Laurea Magistrale in Medical Bioinformatics*

*Nicola Bombieri – Federico Busato*

# Agenda

- **Introduction to OpenMP**
  - What is OpenMP?
  - OpenMP execution and memory model
- **OpenMP basic syntax**
  - Parallel regions
  - Loops
  - Data sharing
- **Examples and exercises**
- **References**

# What is OpenMP?

*“OpenMP (Open Multi-Processing) is an Application Programming Interface (API) that supports multi-platform shared memory parallel programming in C/C++ and Fortran on many architectures, including Unix and Microsoft Windows platforms”*

- It consists of a set of compiler directives, library routines, and environment variables that influence run-time behavior

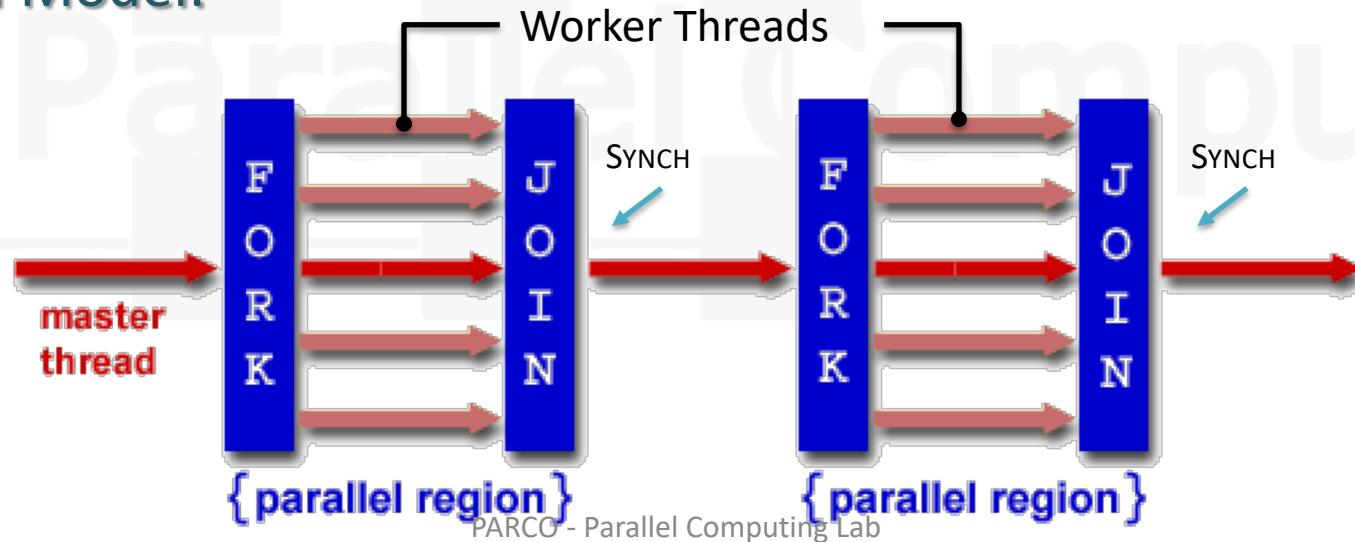
# Creating an OpenMP program (from a sequential one)

- 1) Identify the parallelism that sequential program contains (i.e. finding instructions or regions of code that may be executed concurrently by different processors)
- 2) Express the parallelism that has been identified by using **OPENMP DIRECTIVES**
  - Modifications to the original source program are often needed in just a few places
  - In general, the main effort in parallelizing a program with OpenMP lies in identifying the parallelism (step 1)

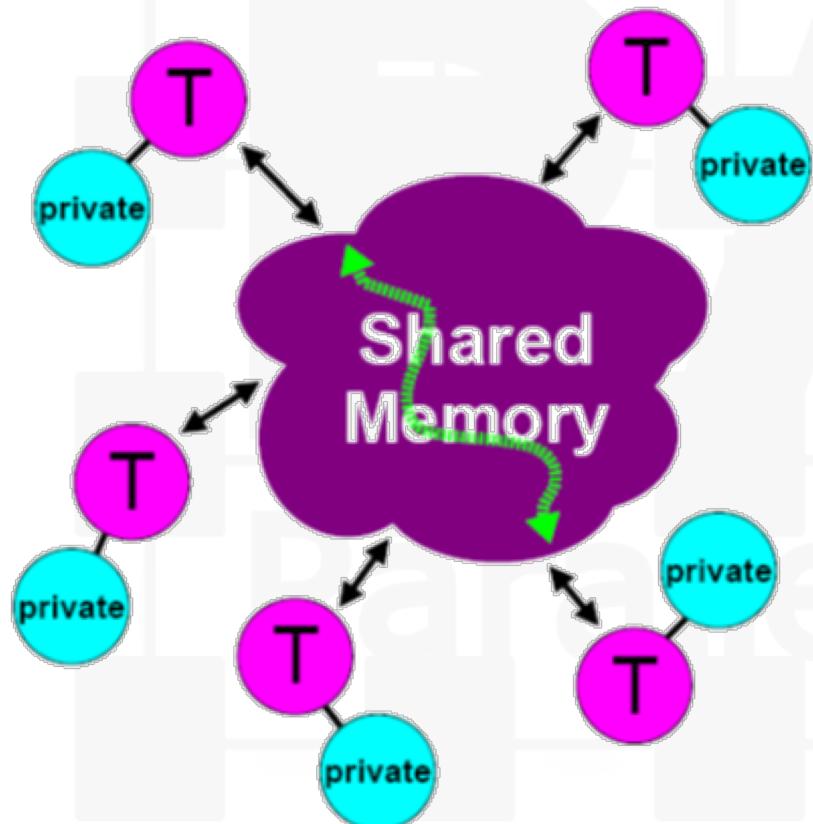
# The OpenMP Execution Model

- Execution starts with one thread of control (master thread)
- Parallel regions fork off new threads on entry (team threads)
- Threads join back together at the end of the region (synchronization)

## Fork-Join Model:



# The OpenMP Memory Model



- All threads have access to the same globally shared memory
- Data can be shared or private
  - Shared data is accessible by all threads
  - Private data can only be accessed by the thread that owns it

# The OpenMP Terminology

- A *parallel region* is a block of code executed by all threads simultaneously
- An OpenMP *team* is the group of threads that currently execute
- A *work-sharing construct* divides the execution of the current code region among the members of a team (i.e. it splits work among threads)

# The OpenMP Syntax

- All OpenMP directives in C/C++ are indicated with a **#pragma omp** and are case sensitive

```
#pragma omp <directive> [clause [clause] ...]
```

- A **#pragma** directive is the method specified by the C standard for providing additional information to the compiler
- Many OpenMP directives are followed by clauses to specify additional information
- The pragma usually applies only to the statement (or block of code) immediately following it

# The PARALLEL Pragma (I)

- The **PARALLEL** pragma starts a parallel region
- It creates a team of N threads (where N is determined at runtime, usually from the number of CPU cores), all of which execute the next statement (or the next block, if the statement is enclosed in {...})
- After the statement, the threads join back into one

```
#pragma omp parallel
{
    // Code inside this region runs in parallel
    printf("Hello!\n");
}
```

- This example prints the text "Hello!" followed by a newline as

```
$ Hello!
$ Hello!
$ Hello!
$ Hello!
```

# The PARALLEL Pragma (II)

- You can explicitly specify the number of threads to be created in the team using the `NUM_THREADS` clause:

```
#pragma omp parallel num_threads(<N>)
```

```
#pragma omp parallel num_threads(3)
{
    // Code inside this region runs in parallel
    printf("Hello!\n");
}
```

- E.g.:

```
$ Hello!
$ Hello!
$ Hello!
```

# The FOR Pragma (I)

- The `for` directive splits the for-loop so that each thread in the current team handles a different portion of the loop

```
#pragma omp parallel for
for (int n = 0; n < 10; ++n)
    printf("%d ", n);
```

- It is a work-sharing construct, so it does not launch any new thread  
→ it must be enclosed in a parallel region
- There is a shorthand to combine both these directives:

# The FOR Pragma (II): an example

```
#pragma omp parallel for
{
    for (int i = 0; i < 1024; ++i)
        C[i] = A[i] + B[i];
}
```

Thread 0	Thread 1	Thread 2	Thread 3
$i \in [0 - 255]$	$i \in [256 - 511]$	$i \in [512 - 767]$	$i \in [768 - 1023]$
A[i]	A[i]	A[i]	A[i]
+	+	+	+
B[i]	B[i]	B[i]	B[i]
=	=	=	=
C[i]	C[i]	C[i]	C[i]

# The FOR ORDERED Pragma

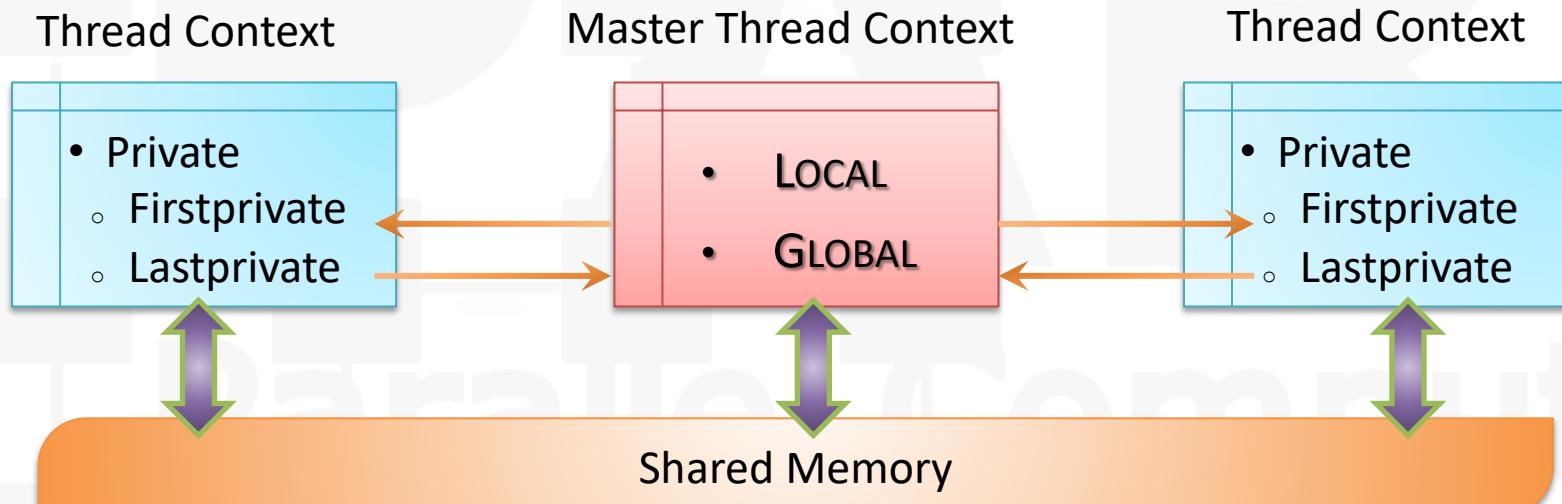
- The order in which loop iterations are executed is unspecified, it depends on runtime conditions
- However it is possible to force a block within a loop to be executed in the order in which loop iterations would be sequentially executed

```
#pragma omp for ordered
for (int n = 0; n < 10; ++n) {
    ...
    #pragma omp ordered
    ...
}
```

- Specifies that code under a parallelized for loop should be executed like a sequential loop
- There may be only one ordered block in the enclosing loop, which must contain the ordered clause

# Data sharing (I)

Variable Scope:



# Data sharing (II)

## The **SHARED** Pragma

- The shared clause indicates that each variable in its list is shared, i.e. all threads access the same variable

```
#pragma omp parallel shared(<var1>, <var2>, ... )
```

- A shared variable exists in only one memory location and all threads can read or write to that address
- It is the programmer's responsibility to ensure that multiple threads properly access SHARED variables (such as via CRITICAL sections)
- Because OpenMP is based upon the shared memory programming model, most variables are shared by default

# Data sharing (III)

## The PRIVATE Pragma

- In a parallel region it is possible to specify which variables are shared between threads and which are not
- The **private** clause indicates that each variable in its list is *private*, i.e. each thread has its own copy of it

```
#pragma omp parallel private(<var1>, <var2>, ... )
```

- Private variables are undefined on entry and exit of the parallel region, i.e. a **private copy** is an uninitialized variable having the same name and the same type as the original variable
- The following are ***implicit private variables***:
  - Loop index variables
  - Stack variables in subroutines called from parallel regions

# Data sharing (IV)

## The FIRSTPRIVATE and LASTPRIVATE Pragma

- A private variable within a parallel region **does not** copy the value of the same variable **inside** and **outside** of the region
- 

- To override this behavior, use the FIRSTPRIVATE clause instead

```
#pragma omp parallel firstprivate(<var1>, <var2>, ... )
```

- All private copies of variables in the list are initialized with the value of the corresponding original variable before entering the region
- 

- To preserve the last value of a private variable on exit of the parallel region, use the LASTPRIVATE clause

```
#pragma omp parallel lastprivate(<var1>, <var2>, ... )
```

- All original variables corresponding to those in the list are updated after the end of the parallel region

# RunTime Functions (I)

- Include `omp` Header

```
#include <omp.h>
```

- OpenMP C/C++ application program interface (API) for threads:
  - `int omp_get_max_threads()`  
Returns the maximum number of threads that would be available in a parallel region
  - `int omp_get_num_threads()`  
Returns the number of threads in the parallel region
  - `int omp_get_thread_num()`  
Returns the thread number of the thread executing within its thread team
  - `void omp_set_num_threads(int num_threads)`  
Sets the number of threads in subsequent parallel regions, unless overridden by a `num_threads` clause.

# RunTime Functions (II)

- OpenMP C/C++ application program interface (API):
  - `double omp_get_wtime()`  
Returns a value in seconds of the time elapsed from some point.
  - `void omp_set_dynamic(int val)`  
Indicates that the number of threads available in subsequent parallel region can be adjusted by the run time
  - `int omp_in_parallel()`  
Returns nonzero if called from within a parallel region
  - `int omp_get_num_procs()`  
Returns the number of processors that are available when the function is called

# Examples

- To compile a program containing OpenMP directives you have to use gcc/g++ and add the **-fopenmp** flag:

```
g++ -fopenmp source.cpp -o out.x
```

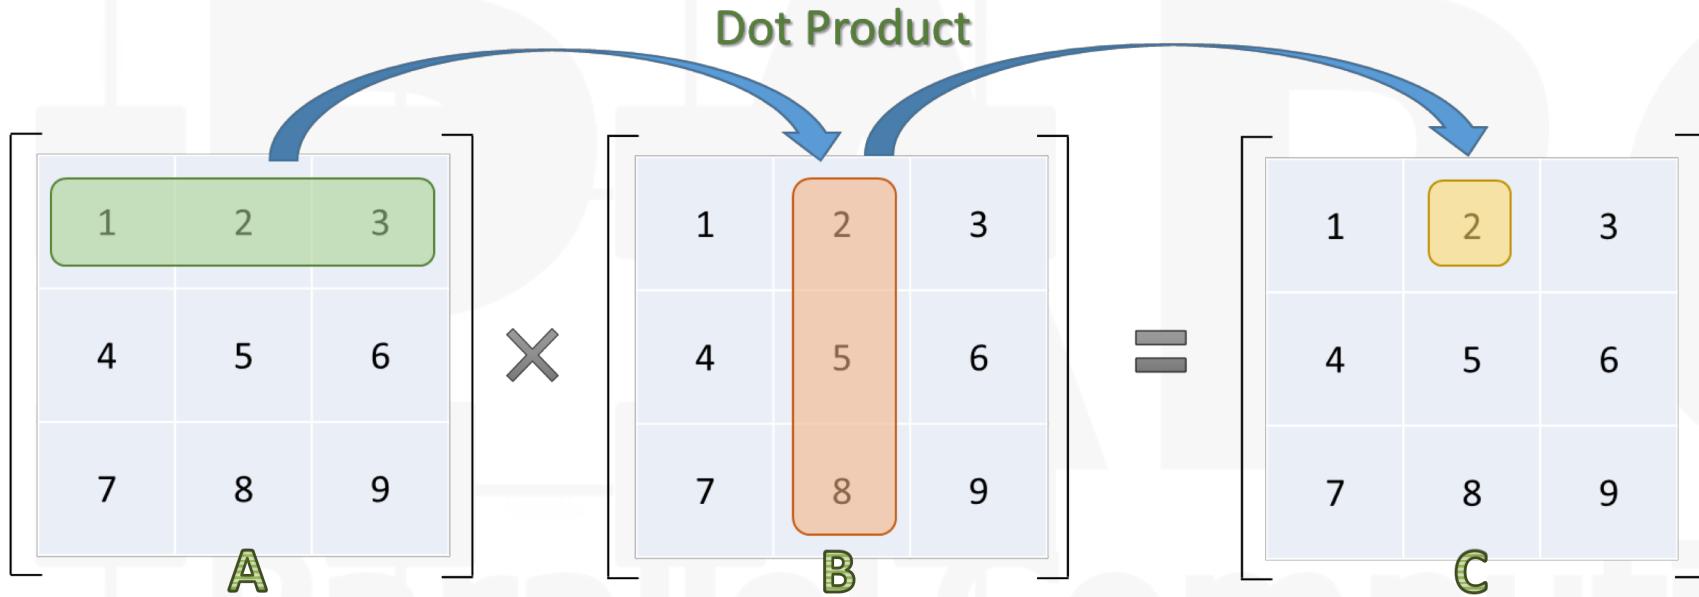
- gcc/g++ OpenMP Support:
  - GCC 4.2 implements version 2.5
  - GCC 4.4 implements version 3.0
  - GCC 4.7 implements version 3.1
  - GCC 4.9 implements version 4.0
  - GCC 5.x implements device Offloading

- Examples:**
  - Hello World
  - Test Runtime Functions
  - Test PRIVATE / FIRSTPRIVATE / LASTPRIVATE Directives



# Exercises (I)

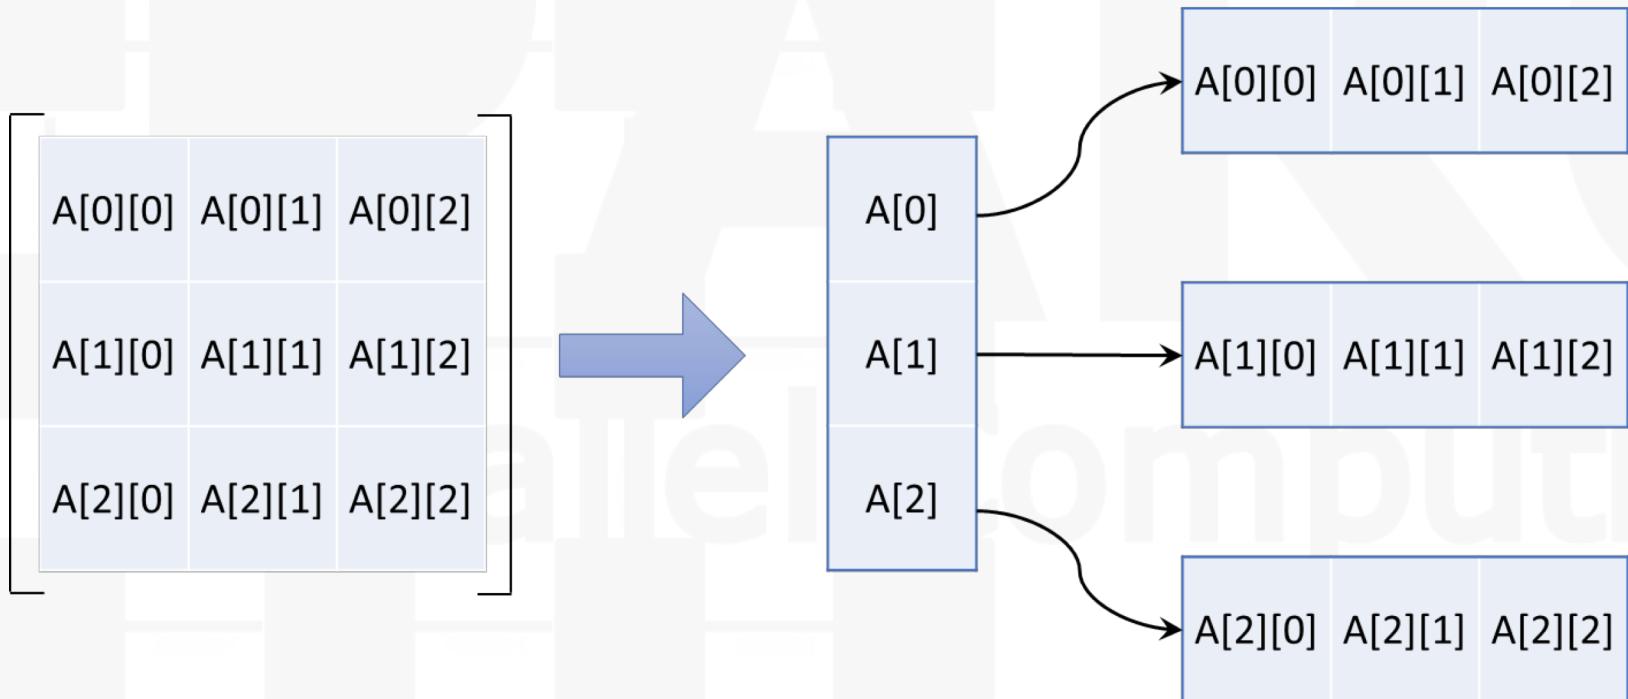
## MATRIX MULTIPLICATION



```
for row = 0 to N
    for col = 0 to N
        sum = 0
        for k = 0 to N
            sum += A[row][k] * B[k][col]
        C[row][col] = sum
    endfor
endfor
```

# Exercises (I)

## MATRIX MULTIPLICATION



# Exercises (II)

## Simple exercises:

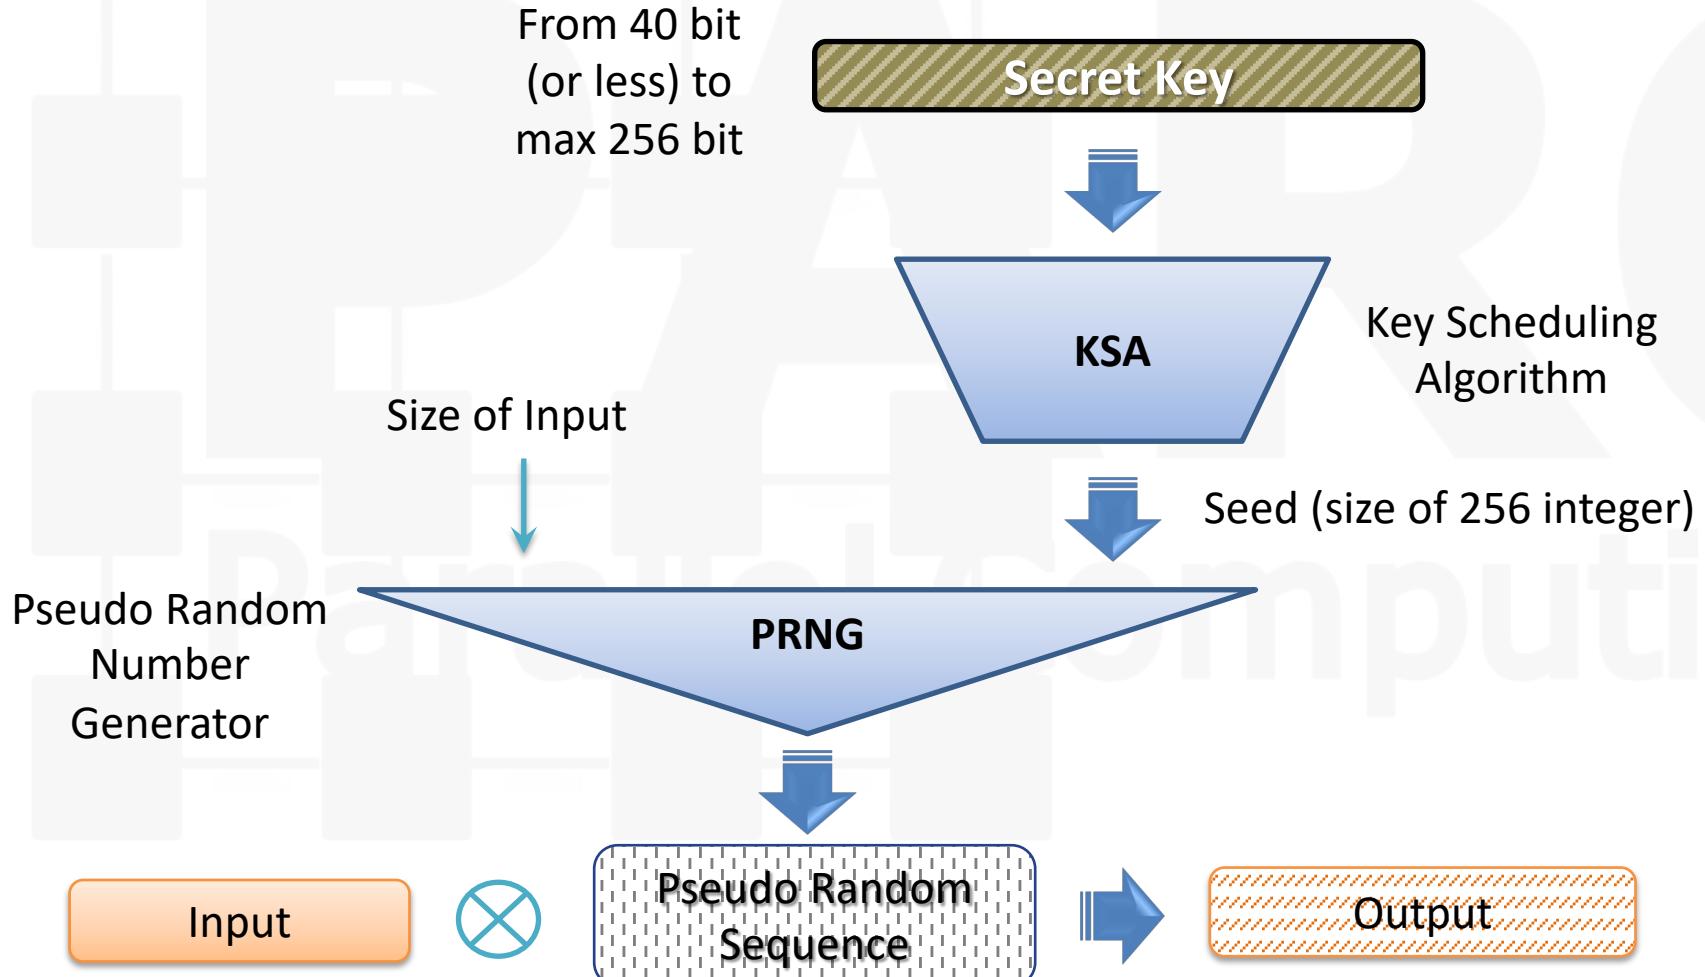
- Factorial( $n$ )
  - $n!$
- Find:
  - Find two (given) consecutive numbers in an array

## Little more complex:

- RC4 Cipher



# RC4 Cipher



# References

- Official OpenMP Specification (v4.0.0)  
[www.openmp.org/mp-documents/OpenMP4.0.0.pdf](http://www.openmp.org/mp-documents/OpenMP4.0.0.pdf)
- OpenMP in Visual C++  
<https://msdn.microsoft.com/en-us/library/tt15eb9t.aspx>
- OpenMP Tutorial - Livermore Computing Center  
<https://computing.llnl.gov/tutorials/openMP/>
- Guide into OpenMP: Easy multithreading programming for C++  
<http://bisqwit.iki.fi/story/howto/openmp/>
- A “Hands-on” Introduction to OpenMP  
<http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>